

## HW6

```
#1a
set.seed(3301)
dat1 = read.table("trees.txt")

dat= dat1[sample(1:nrow(dat1)),]
dat

##      D  H   V
## 10 11.2 75 19.9
## 24 16.0 72 38.3
##  9 11.1 80 22.6
## 25 16.3 77 42.6
## 14 11.7 69 21.3
## 29 18.0 80 51.5
## 16 12.9 74 22.2
##  3  8.8 63 10.2
## 22 14.2 80 31.7
## 26 17.3 81 55.4
##  1  8.3 70 10.3
## 23 14.5 74 36.3
##  2  8.6 65 10.3
## 21 14.0 78 34.5
## 20 13.8 64 24.9
##  5 10.7 81 18.8
##  8 11.0 75 18.2
## 11 11.3 79 24.2
## 31 20.6 87 77.0
## 18 13.3 86 27.4
##  4 10.5 72 16.4
## 28 17.9 80 58.3
## 12 11.4 76 21.0
## 19 13.7 71 25.7
## 13 11.4 76 21.4
## 17 12.9 85 33.8
##  6 10.8 83 19.7
## 27 17.5 82 55.7
## 30 18.0 80 51.0
##  7 11.0 66 15.6
## 15 12.0 75 19.1

X1 = cbind(1, dat$D, (dat$D)^2, dat$D*dat$H, dat$H, (dat$H)^2)
y1 = log(dat$V)

olscv=function(X, y, K=5, permute=FALSE)
{
```

```

n=length(y)
if(permute)
{
ind=sample(n)
} else
{
ind=1:n
}
total.sq.err=0
for(k in 1:K)
{
leave.out=ind[ (1+floor((k-1)*n/K)):floor(k*n/K) ]
X.tr=X[-leave.out,,drop=FALSE]
y.tr=y[-leave.out]
X.va=X[leave.out,,drop=FALSE]
y.va=y[leave.out]
bhat.tr=lm.fit(x=X.tr, y=y.tr)$coefficients
total.sq.err=total.sq.err + sum((y.va - X.va%*%bhat.tr)^2 )
}
return(total.sq.err/n)
}
get.ic=function(X, y, K=5)
{
n=dim(X)[1]
p=dim(X)[2]
beta.hat=lm.fit(x=X, y=y)$coefficients
rss=sum((y-X%*%beta.hat)^2)
common=n*log(2*pi)+n*log(rss/n) + n
aic=common + 2*(p+1)
bic=common + (p+1)*log(n)
est.mspe=olscv(X=X,y=y, K=K)
return(c(aic, bic, est.mspe))
}
keep.status=as.matrix(expand.grid( replicate(5, c(0,1), simplify=FALSE) ))
scores=matrix(NA, nrow=nrow(keep.status), ncol=3)
for(j in 1:nrow(keep.status))
{
keep=as.logical(c(1, keep.status[j,]))
scores[j,] = get.ic(X=X1[,keep, drop=FALSE], y=y1, K=5)
}
result=cbind(scores, keep.status)
colnames(result)=c("AIC", "BIC", "Est. MSPE", "D", "D^2", "D*H", "H", "H^2")
result

```

```

##           AIC           BIC   Est. MSPE D D^2 D*H H H^2
## [1,]  51.15694  54.02492 0.283139819 0  0  0 0 0
## [2,] -33.90673 -29.60477 0.019028369 1  0  0 0 0
## [3,] -18.80140 -14.49943 0.032889688 0  1  0 0 0
## [4,] -38.88103 -33.14508 0.015980573 1  1  0 0 0
## [5,] -38.92627 -34.62431 0.018611391 0  0  1 0 0
## [6,] -43.10662 -37.37067 0.016620345 1  0  1 0 0
## [7,] -37.12530 -31.38935 0.019329150 0  1  1 0 0
## [8,] -59.96195 -52.79202 0.007443704 1  1  1 0 0
## [9,]  36.25653  40.55849 0.171288264 0  0  0 1 0

```

```
## [10,] -51.98466 -46.24871 0.011119744 1 0 0 1 0
## [11,] -31.42541 -25.68946 0.023214911 0 1 0 1 0
## [12,] -61.31366 -54.14373 0.006938839 1 1 0 1 0
## [13,] -38.02440 -32.28845 0.018886283 0 0 1 1 0
## [14,] -59.59979 -52.42985 0.007137994 1 0 1 1 0
## [15,] -42.02996 -34.86003 0.015004631 0 1 1 1 0
## [16,] -59.46716 -50.86324 0.007396804 1 1 1 1 0
## [17,] 36.38881 40.69077 0.171552384 0 0 0 0 1
## [18,] -50.86370 -45.12775 0.011854996 1 0 0 0 1
## [19,] -30.34839 -24.61244 0.024502576 0 1 0 0 1
## [20,] -60.82499 -53.65505 0.007112102 1 1 0 0 1
## [21,] -38.38371 -32.64776 0.018856739 0 0 1 0 1
## [22,] -58.40831 -51.23838 0.007673977 1 0 1 0 1
## [23,] -45.57566 -38.40573 0.013220505 0 1 1 0 1
## [24,] -58.86482 -50.26090 0.007337706 1 1 1 0 1
## [25,] 38.22342 43.95937 0.184654972 0 0 0 1 1
## [26,] -53.42026 -46.25033 0.011878035 1 0 0 1 1
## [27,] -34.78626 -27.61633 0.022726751 0 1 0 1 1
## [28,] -59.74039 -51.13647 0.008193491 1 1 0 1 1
## [29,] -42.46605 -35.29612 0.018378160 0 0 1 1 1
## [30,] -57.67335 -49.06943 0.008251693 1 0 1 1 1
## [31,] -46.99104 -38.38711 0.014417017 0 1 1 1 1
## [32,] -57.75516 -47.71725 0.008321643 1 1 1 1 1
```

```
-61.31366 -54.14373 0.006938839 1 1 0 1 0 subset("D", "D^2", H)
```

```
#2a
```

```
df1 = read.table("divorce.txt")
```

```
dfs= df1[sample(1:nrow(df1)),]
```

```
X = cbind(1, dfs$year, dfs$unemployed, dfs$femlab, dfs$marriage, dfs$birth, dfs$military)
```

```
y=dfs$divorce
```

```
keep.status=as.matrix(expand.grid(replicate(6, c(0,1), simplify=FALSE) ))
```

```
scores=matrix(NA, nrow=nrow(keep.status), ncol=3)
```

```
for(j in 1:nrow(keep.status))
```

```
{
```

```
keep=as.logical(c(1, keep.status[j,]))
```

```
scores[j,] = get.ic(X=X[,keep,drop=FALSE], y=y, K=5)
```

```
}
```

```
result=cbind(scores, keep.status)
```

```
colnames(result)=c("AIC", "BIC", "Est. MSPE", "year", "unemployed", "femlab", "marriage", "birth", "mil.
```

```
result
```

```
##          AIC      BIC Est. MSPE year unemployed femlab marriage birth
## [1,] 488.7042 493.3918 32.683423 0 0 0 0 0
## [2,] 376.5083 383.5397 7.301734 1 0 0 0 0
## [3,] 487.2109 494.2423 31.469088 0 1 0 0 0
## [4,] 378.5012 387.8764 7.317050 1 1 0 0 0
## [5,] 354.7949 361.8263 5.505102 0 0 1 0 0
## [6,] 351.1220 360.4972 5.189291 1 0 1 0 0
## [7,] 356.5218 365.8970 5.507413 0 1 1 0 0
## [8,] 352.4142 364.1332 5.206033 1 1 1 0 0
## [9,] 464.8266 471.8580 25.301632 0 0 0 1 0
## [10,] 378.4685 387.8437 7.545192 1 0 0 1 0
## [11,] 450.5445 459.9198 20.511985 0 1 0 1 0
## [12,] 380.4676 392.1867 7.708269 1 1 0 1 0
```

##	[13,]	354.2976	363.6729	5.605174	0	0	1	1	0
##	[14,]	348.6569	360.3759	5.264765	1	0	1	1	0
##	[15,]	352.8331	364.5521	5.425902	0	1	1	1	0
##	[16,]	342.0435	356.1063	4.786245	1	1	1	1	0
##	[17,]	433.9999	441.0313	15.913895	0	0	0	0	1
##	[18,]	350.4289	359.8041	5.300500	1	0	0	0	1
##	[19,]	390.9631	400.3383	8.875145	0	1	0	0	1
##	[20,]	333.3542	345.0732	4.120880	1	1	0	0	1
##	[21,]	332.3504	341.7256	4.178502	0	0	1	0	1
##	[22,]	330.3148	342.0338	4.036506	1	0	1	0	1
##	[23,]	322.1993	333.9183	3.587583	0	1	1	0	1
##	[24,]	322.6855	336.7483	3.627496	1	1	1	0	1
##	[25,]	435.3220	444.6972	17.112431	0	0	0	1	1
##	[26,]	338.3458	350.0648	4.724621	1	0	0	1	1
##	[27,]	389.7720	401.4910	9.375392	0	1	0	1	1
##	[28,]	328.6932	342.7560	4.144707	1	1	0	1	1
##	[29,]	305.7125	317.4315	3.144683	0	0	1	1	1
##	[30,]	297.2075	311.2704	2.860362	1	0	1	1	1
##	[31,]	305.1557	319.2186	3.080686	0	1	1	1	1
##	[32,]	298.9332	315.3399	2.883705	1	1	1	1	1
##	[33,]	490.6776	497.7090	32.940527	0	0	0	0	0
##	[34,]	378.4579	387.8331	7.406861	1	0	0	0	0
##	[35,]	488.7953	498.1705	31.613396	0	1	0	0	0
##	[36,]	380.4578	392.1768	7.468437	1	1	0	0	0
##	[37,]	356.4381	365.8133	5.593542	0	0	1	0	0
##	[38,]	351.4301	363.1492	5.292228	1	0	1	0	0
##	[39,]	358.3422	370.0613	5.646839	0	1	1	0	0
##	[40,]	353.2910	367.3538	5.365377	1	1	1	0	0
##	[41,]	463.9449	473.3201	25.097068	0	0	0	1	0
##	[42,]	380.4404	392.1594	7.756714	1	0	0	1	0
##	[43,]	452.3988	464.1179	21.152696	0	1	0	1	0
##	[44,]	382.4339	396.4967	7.889235	1	1	0	1	0
##	[45,]	354.5405	366.2596	5.674388	0	0	1	1	0
##	[46,]	343.7198	357.7826	5.086314	1	0	1	1	0
##	[47,]	353.9535	368.0163	5.534376	0	1	1	1	0
##	[48,]	338.5314	354.9380	4.679825	1	1	1	1	0
##	[49,]	433.5884	442.9636	15.829247	0	0	0	0	1
##	[50,]	350.6562	362.3752	5.204449	1	0	0	0	1
##	[51,]	392.0134	403.7324	8.981115	0	1	0	0	1
##	[52,]	335.2648	349.3276	4.200108	1	1	0	0	1
##	[53,]	334.1290	345.8481	4.240597	0	0	1	0	1
##	[54,]	332.2960	346.3588	4.189880	1	0	1	0	1
##	[55,]	323.5488	337.6116	3.670322	0	1	1	0	1
##	[56,]	323.4006	339.8073	3.715725	1	1	1	0	1
##	[57,]	434.0992	445.8183	17.114144	0	0	0	1	1
##	[58,]	340.2344	354.2973	4.790949	1	0	0	1	1
##	[59,]	391.3723	405.4351	9.576873	0	1	0	1	1
##	[60,]	330.0888	346.4954	4.180049	1	1	0	1	1
##	[61,]	305.6053	319.6682	3.122805	0	0	1	1	1
##	[62,]	289.8468	306.2534	2.599032	1	0	1	1	1
##	[63,]	303.4288	319.8355	3.010537	0	1	1	1	1
##	[64,]	290.9269	309.6773	2.614563	1	1	1	1	1
##	military								
##	[1,]	0							

##	[2,]	0
##	[3,]	0
##	[4,]	0
##	[5,]	0
##	[6,]	0
##	[7,]	0
##	[8,]	0
##	[9,]	0
##	[10,]	0
##	[11,]	0
##	[12,]	0
##	[13,]	0
##	[14,]	0
##	[15,]	0
##	[16,]	0
##	[17,]	0
##	[18,]	0
##	[19,]	0
##	[20,]	0
##	[21,]	0
##	[22,]	0
##	[23,]	0
##	[24,]	0
##	[25,]	0
##	[26,]	0
##	[27,]	0
##	[28,]	0
##	[29,]	0
##	[30,]	0
##	[31,]	0
##	[32,]	0
##	[33,]	1
##	[34,]	1
##	[35,]	1
##	[36,]	1
##	[37,]	1
##	[38,]	1
##	[39,]	1
##	[40,]	1
##	[41,]	1
##	[42,]	1
##	[43,]	1
##	[44,]	1
##	[45,]	1
##	[46,]	1
##	[47,]	1
##	[48,]	1
##	[49,]	1
##	[50,]	1
##	[51,]	1
##	[52,]	1
##	[53,]	1
##	[54,]	1
##	[55,]	1

```
## [56,]      1
## [57,]      1
## [58,]      1
## [59,]      1
## [60,]      1
## [61,]      1
## [62,]      1
## [63,]      1
## [64,]      1
```

```
subset(year, femlab, marriage, birth, military) 289.8468 306.2534 2.496206 1 0 1 1 1 1
```

```
#2b
#model with predictors that generated the lowest AIC, BIC, MSPE
train=1:70
test= 71:77
tdata= df1[sample(1:70),]
ttdata = df1[sample(test),]
y=tdata$divorce
#{year, unemployed, femlab, marriage, birth, military}
X = cbind(1, tdata$year, tdata$unemployed, tdata$femlab, tdata$marriage, tdata$birth, tdata$military)
X1 = cbind(1, ttdata$year, ttdata$unemployed, ttdata$femlab, ttdata$marriage,
           ttdata$birth, ttdata$military)
y1= ttdata$divorce
get.ic(X=X[train,], y=y[train], K=5)
```

```
## [1] 262.112699 280.100661 2.875636
```

```
#{year, femlab, marriage, birth, military}
X2 = cbind(1, tdata$year, tdata$femlab, tdata$marriage, tdata$birth, tdata$military)
X3 = cbind(1, ttdata$year, ttdata$femlab, ttdata$marriage,
           ttdata$birth, ttdata$military)
get.ic(X=X2[train,], y=y[train], K=5)
```

```
## [1] 261.231001 276.970468 2.900803
```

According to the output the subset of predictors with the best output are {year, femlab, marriage, birth, military}

```
#2c
#FULL model
beta.hat1=qr.coef(qr(X[train,]), y=y[train])
beta.hat1
```

```
## [1] 428.16156307 -0.22860248 -0.05325023 0.87878246 0.13802516
## [6] -0.11343592 -0.04909694
```

```
fitted1 = X1[1:7,]%*%beta.hat1 #X1 is the test set
mean((y1[1:7] - fitted1)^2)
```

```
## [1] 5.112417
```

```
#{year, femlab, marriage, birth, military}
beta.hat2=qr.coef(qr(X2[train,]), y=y[train])
beta.hat2
```

```
## [1] 456.54703307 -0.24504472 0.93050659 0.14834382 -0.10601865
## [6] -0.04747228
```

```
fitted2 = X3[1:7,]%*%beta.hat2 #X3 is the test set
residuals2 = y[test] - fitted2
mean((y1[1:7] - fitted2)^2)
```

```
## [1] 5.095095
```

yes, the model does overestimate the response for subjects in the test set because the values of the mean squared error is greater than 0

```
#3a
data = read.table("paper.txt")
data
```

```
##      bright operator
## 1      59.8        a
## 2      60.0        a
## 3      60.8        a
## 4      60.8        a
## 5      59.8        a
## 6      59.8        b
## 7      60.2        b
## 8      60.4        b
## 9      59.9        b
## 10     60.0        b
## 11     60.7        c
## 12     60.7        c
## 13     60.5        c
## 14     60.9        c
## 15     60.3        c
## 16     61.0        d
## 17     60.8        d
## 18     60.6        d
## 19     60.5        d
## 20     60.5        d
```

```
X = cbind(rep(1,20), 1*(data$operator=='a'), 1*(data$operator=='b'), 1*(data$operator=='c')) #full mode
X0 = cbind(rep(1,20)) #null model with just the intercept
print("assumptions: intercept, predictor variable, coefficients, and the response as bright")
```

```
## [1] "assumptions: intercept, predictor variable, coefficients, and the response as bright"
```

```
y = data$bright
beta.hat = qr.coef(qr(X), y=y)
beta.hat
```

```
## [1] 60.68 -0.44 -0.62 -0.06
```

```
y
```

```
## [1] 59.8 60.0 60.8 60.8 59.8 59.8 60.2 60.4 59.9 60.0 60.7 60.7 60.5 60.9 60.3
## [16] 61.0 60.8 60.6 60.5 60.5
```

```
rssf= sum((y-X%*%beta.hat)^2)
beta.hat0= qr.coef(qr(X0), y=y)
rssf0 = sum((y-X0%*%beta.hat0)^2)
n = length(y); p = length(beta.hat); d=4
f=((rssf0-rssf)/3)/(rssf/(n-p))
1-pf(f,d,n-p)
```

```
## [1] 0.01628656
```

$H_0: B_2=B_3=B_4=0$   $H_1: H_0$  is false  $B_1$  is the intercept pvalue  $>.01$  we do not reject the null hypothesis at .01 significance level and conclude that the operator is not relevant and significant at .01 in the linear regression model with response bright.

```
#3b
```

```
X = cbind(1,data$operator)
y = data$bright
get.ic(X=X, y=y)
```

```
## [1] 18.2116991 21.1988960 0.1748935
```

```
X1= cbind(rep(1,20))
get.ic(X=X1, y=y)
```

```
## [1] 23.0800462 25.0715107 0.1797031
```

disregard the last 3rd column, which is est.mse

According to the values of the AIC and BIC, the model with the predictors are more accurate than the model with only the intercept. However, the values of the AIC and BIC also reveals that the difference is between the two models are not that significant. With the AIC model being 18.2117 for the model with predictor and 23.08005 for the model with only the intercept. And the BIC model 21.1989 and 25.07151. This demonstrates that the conclusion we had in part a is feasible due to the lack of significance of the predictor variable, operator.

```
#4
```

```
set.seed(3301)
```

```
get.bic=function(X, y)
{
  n=dim(X)[1]
  p=dim(X)[2]
  beta.hat=lm.fit(x=X, y=y)$coefficients
  rss=sum((y-X%%beta.hat)^2)
  bic=n*log(2*pi)+n*log(rss/n) + n + log(n)*(p+1)
  return(bic)
}
```

```
reps=5e4
beta=c(1, 0, 0)
sigma=0.5
n=c(5,10,20,50,100)
m.list = numeric(5)
n.list=numeric(reps)
```

```
X1 = cbind(rep(1,100),0,0)
for(k in 1:100){
  for(s in 2:3){
    X1[k,s]= rnorm(1)
  }
}
for(i in 1:5){
  X = X1[1:n[i],]
}
for(r in 1:reps){
```



```

y=X%*%beta + sigma*rnorm(n=n[i])
bic.list=numeric(4)
bic.list[1]=get.bic(X=X, y=y)
bic.list[2]=get.bic(X=X[, -3], y=y)
bic.list[3]=get.bic(X=X[, -2], y=y)
bic.list[4]= get.bic(X=cbind(rep(1,n[i])), y=y)
picked.index= which.min(bic.list)
n.list[r]=1*(picked.index==1)
}
m.list[i] = mean(n.list)
}
m.list

```

```
## [1] 0.27006 0.06098 0.01608 0.00404 0.00152
```